

LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Chintan Maneesh Gangurde
Roll Number	MHM2026003
Programme / Department	M.Tech in IT (HMIIT)
Topic ID	T_8
Full Assigned Topic	Error Propagation in Multi-Step Agentic AI Workflows for Research Automation
AI Tool Used	ChatGPT
Date	21 August 2026

IMPORTANT: You will receive a topic from T1-T40. Your task is not to produce the “best” paper. Your task is to generate an AI-written research paper and systematically determine whether its citations can be trusted.

1. PURPOSE OF THIS LAB

Generative AI systems can produce research papers that look scholarly and may contain convincing references, author names, journal titles, DOIs, arXiv IDs, and in-text citations. However, the presence of a citation does not guarantee that the cited publication actually exists, that its metadata are correct, that the identifier belongs to that publication, or that the cited paper supports the statement made by the AI.

In this laboratory activity, you will test these issues yourself by generating a research paper and auditing a systematic sample of its references.

Question 1 - Citation Authenticity: Does the cited publication actually exist, and is it correctly identified?

Question 2 - Claim-Citation Support: If the publication is genuine, does it actually support the claim for which the AI cited it?

2. IMPORTANT RULES

- Use only the topic assigned to you by the instructor.
- Generate the paper in a fresh AI conversation.
- Do not tell the AI that you are conducting a citation audit or hallucination test.
- Do not intentionally ask the AI to generate fake references.
- Preserve the original AI-generated paper exactly as produced.
- Do not correct, regenerate, replace, or question references before saving the original output.
- For verification, do not accept another AI chatbot's statement as proof that a reference exists.
- Use scholarly databases, publisher pages, DOI/arXiv/PMID records, or other reliable scholarly sources.
- Record what you actually find, even if all references are genuine or all are problematic.

Most important principle: You are not being evaluated on whether the AI produces good or bad references. You are being evaluated on how systematically and correctly you audit them.

3. WHAT YOU WILL DO

Part	Activity
------	----------

A	Record your baseline perception of AI citations
B	Record the AI environment used
C	Generate and preserve an AI-written research paper
D	Create a reference inventory
E	Select 10 references systematically
F	Judge their plausibility before verification
G	Verify authenticity, metadata and identifiers
H	Determine whether genuine references support the AI-generated claims

4. RECOMMENDED 2-HOUR TIME PLAN

Time	Activity
0-10 min	Parts A-B: Baseline and AI environment
10-30 min	Part C: Generate and save research paper
30-40 min	Parts D-E: Inventory and select references
40-45 min	Part F: Pre-verification plausibility
45-80 min	Part G: Authenticity and metadata audit
80-110 min	Part H: Claim-citation entailment audit
110-120 min	Calculate scores, check worksheet and prepare submission

Do not spend excessive time formatting the AI-generated paper. The audit is the main activity.

PART A - BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

Never Rarely Occasionally Weekly ☒ Almost daily

A2. Have you previously asked an AI system to generate references?

Frequently ☒ Occasionally Once or twice Never

A3. Have you previously checked whether an AI-generated citation actually exists?

☒ Always Sometimes Rarely Never Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

☒ 1 2 3 4 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B - RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	ChatGPT
Model name	GPT-5.6 Luna
Model/version shown	GPT-5.6 Luna
Access tier	☒ Free Paid/Pro Institutional Not sure Selected: Free
Web browsing available?	☒ Yes No Not sure Selected: Yes
Was web browsing actually used?	☒ Yes No Not sure Selected: Yes
Research/Deep Research mode used?	Yes ☒ No

	Selected: No
Date/time generated	21 August 2026, 22:06 IST

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C - GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

Generate a scholarly research paper on the topic: "[YOUR ASSIGNED TOPIC]". The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.

C3. Preserve the Original Output

Save the complete AI-generated paper

Save the complete reference list

Take screenshot(s) showing the generated output/references

Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

AI-GENERATED RESEARCH PAPER

Error Propagation in Multi-Step Agentic AI Workflows for Research Automation

Abstract

Agentic AI systems are increasingly being used for research tasks that involve several connected steps, such as searching for information, selecting sources, extracting evidence, summarizing findings, and writing a final report. Systems such as ReAct, Toolformer and AutoGen show how language models can combine reasoning, tools and multiple agents. However, a multi-step workflow also creates a reliability problem: an error made early in the process can be passed to later steps and become harder to notice. This paper discusses how such errors can propagate through retrieval, source selection, extraction, planning, memory, communication and final writing. It also reviews approaches such as external grounding, tool use, reflection and independent verification. The main argument is that reliable research automation should treat intermediate results as uncertain evidence and should check important claims against their original sources.

Keywords

agentic AI; error propagation; multi-step workflows; research automation; LLM agents; verification; reliability

1. Introduction

Large language models are moving from simple question-answering systems toward agents that can plan tasks, use tools, read external information and take several actions before producing an answer. ReAct combines reasoning and acting so that an agent can interact with an external environment while solving a task [1]. Toolformer studies how language models can decide when to use external APIs and how to incorporate their results [2]. AutoGen provides a framework where multiple agents can communicate and use tools to complete tasks [3].

These capabilities are useful for research automation. A research task can be divided into smaller jobs: understand the question, search for papers, open sources, extract useful evidence, compare findings, create an outline, write a draft, and check the final claims. The difficulty is that the final answer depends on many intermediate outputs. If one step selects the wrong source or misunderstands a result, later steps may use that output as if it were correct.

2. Background and Related Work

ReAct showed that reasoning and external actions can be combined in an interleaved process [1]. Toolformer demonstrated that language models can learn to use external tools such as calculators, search engines and question-answering systems [2]. AutoGen extended the agent idea to multiple communicating agents that can use language, code, tools and human input [3].

Research-focused systems also show the value of external evidence. WebGPT uses web browsing and references to support long-form answers [5]. AgentBench evaluates language models as agents across interactive environments and identifies long-term reasoning and decision-making as important challenges [6]. Reflexion uses verbal feedback and episodic memory to improve later decisions [4], while Self-Refine repeatedly generates feedback and improves an output [7].

3. How Error Propagation Happens

A multi-step research workflow can be viewed as a chain: query → plan → search → source selection → extraction → summary → synthesis → writing → verification. Each stage produces information that can be used by the next stage. If an early result is wrong, a later stage may not know that it is wrong.

Retrieval errors can occur when an agent selects a source that is related to the topic but does not actually answer the question. Extraction errors can happen when a correct paper is found but a result or limitation is misunderstood. Planning errors can send the workflow in an unhelpful direction. Communication errors can occur when one agent passes a summary to another without the original evidence.

Long contexts create another problem. Research on long-context language models found that performance can decrease when relevant information is located in the middle of a long input [8]. This means that giving an agent more documents does not automatically make the final answer more reliable.

4. Error Amplification in Research Automation

Not every error becomes worse. Some errors are corrected by later steps, while others are copied into notes, summaries and final claims. A propagating error is especially dangerous because later agents may treat it as an established fact. Repetition does not make a claim more reliable when every repetition came from the same incorrect intermediate result.

A verification step can break this chain if it returns to the original source. In contrast, a checker that only reads the previous agent's summary may simply repeat the same mistake. Therefore, the independence and quality of verification matter, not just the number of checking steps.

5. Current Approaches to Reducing Errors

External grounding can reduce dependence on model memory by requiring an agent to search for evidence [1,5]. Tool use can provide calculators, search systems and other external functions [2]. Reflection and iterative refinement can improve outputs by asking an agent to inspect and revise its own work [4,7]. Independent verification is especially important for citations, numerical results, dates and central conclusions.

Agent-computer interfaces also matter. SWE-agent shows that a specially designed interface can affect how well an agent performs complex software tasks [9]. The AI Scientist demonstrates how agentic systems can connect idea generation, coding, experiments, visualisation and paper writing into a larger scientific workflow [10]. These systems show both the potential of automation and the importance of checking intermediate results.

6. Research Challenges and Limitations

It is difficult to identify exactly where an error entered a long workflow when many agents and tools are involved. A model can also give a confident explanation even when its evidence is weak. More verification increases time and cost, while self-critique is not fully independent when the same model and assumptions are used for generation and checking. Long documents can also make it harder to use important evidence effectively [8].

7. Research Gaps and Future Directions

A major research gap is the lack of standard methods for measuring error propagation across complete research workflows. Future benchmarks could record intermediate searches, source choices, extracted evidence, agent messages and verification decisions. This would make it easier to identify where a failure started.

Another useful direction is provenance-aware memory. Instead of storing only a summary, an agent could store each claim with its source, location, confidence and date. Later agents could then trace a claim back to the original evidence. Risk-based verification could also be used so that important numerical findings, surprising comparisons and central conclusions receive stronger checks.

8. Conclusion

Multi-step agentic AI can make research faster and more capable, but every additional step creates another place where an error can enter. The main risk is not only that one model makes a mistake; it is that later agents accept the mistake as a valid input and build more work on top of it. Reliable research automation therefore needs source provenance, uncertainty tracking and independent verification. Important claims should be checked against their original sources instead of being trusted because they were repeated by several agents.

References

1. Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing Reasoning and Acting in Language Models. ICLR 2023. arXiv:2210.03629. DOI: 10.48550/arXiv.2210.03629.
2. Schick, T., Dwivedi-Yu, J., Dessi, R., Raileanu, R., Lomeli, M., Hambro, E., Zettlemoyer, L., Cancedda, N., & Scialom, T. (2023). Toolformer: Language Models Can Teach Themselves to Use Tools. NeurIPS 2023. DOI: 10.52202/075280-2997.
3. Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., Jiang, L., Zhang, X., Zhang, S., Awadallah, A., White, R. W., Burger, D., & Wang, C. (2023). AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. arXiv:2308.08155. DOI: 10.48550/arXiv.2308.08155.
4. Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., & Yao, S. (2023). Reflexion: Language Agents with Verbal Reinforcement Learning. arXiv:2303.11366. DOI: 10.48550/arXiv.2303.11366.

5. Nakano, R., Hilton, J., Balaji, S., Wu, J., Ouyang, L., Kim, C., Hesse, C., Jain, S., Kosaraju, V., Saunders, W., Jiang, X., Cobbe, K., Eloundou, T., Krueger, G., Button, K., Knight, M., Chess, B., & Schulman, J. (2021). WebGPT: Browser-assisted question-answering with human feedback. arXiv:2112.09332. DOI: 10.48550/arXiv.2112.09332.
6. Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., Zhang, S., Deng, X., Zeng, A., Du, Z., Zhang, C., Shen, S., Su, Y., Sun, H., Huang, M., Dong, Y., & Tang, J. (2023). AgentBench: Evaluating LLMs as Agents. arXiv:2308.03688. DOI: 10.48550/arXiv.2308.03688.
7. Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhumoye, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck, S., Yazdanbakhsh, A., & Clark, P. (2023). Self-Refine: Iterative Refinement with Self-Feedback. arXiv:2303.17651. DOI: 10.48550/arXiv.2303.17651.
8. Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2023). Lost in the Middle: How Language Models Use Long Contexts. TACL. arXiv:2307.03172. DOI: 10.48550/arXiv.2307.03172.
9. Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S., Narasimhan, K., & Press, O. (2024). SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. NeurIPS 2024. arXiv:2405.15793. DOI: 10.48550/arXiv.2405.15793.
10. Lu, C., Lu, C., Lange, R. T., Foerster, J., Clune, J., & Ha, D. (2024). The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. arXiv:2408.06292. DOI: 10.48550/arXiv.2408.06292.

PART D - PAPER PROFILE AND REFERENCE INVENTORY

Paper title: Error Propagation in Multi-Step Agentic AI Workflows for Research Automation.....

Measure	Value
Approximate pages	About 8 pages including references
Approximate word count	About 2,600 words
Number of sections	8 main sections + Abstract + Keywords + References
Total number of references	10
Approximate in-text citation occurrences	About 20
Did AI warn you to verify references?	No

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference Type	Count
DOI available	10
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	0
URL only	0
No persistent identifier supplied	0
TOTAL	10

Examples

Category	Example	How to Count
DOI available	Kather, J. N., et al. (2019). Predicting survival from colorectal cancer histology slides using deep learning. PLoS Medicine. DOI: 10.1371/journal.pmed.1002730	DOI available
PMID available, no DOI	A biomedical paper for which the AI provides PMID: 12345678, but no DOI is supplied/found	PMID available, no DOI
arXiv ID available	Yao, S., et al. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629	arXiv ID available
URL only	World Health Organization. Colorectal cancer. Available at: https://... with no DOI, PMID or arXiv ID	URL only

No persistent identifier	Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.	No persistent identifier if none is supplied
--------------------------	--	--

Important: Absence of a DOI does not mean a reference is fake. An identifier being present also does not prove authenticity; it must be checked.

PART E - SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Example: If there are 28 references, a reasonable sample could be 1, 2, 3, 8, 13, 18, 23, 26, 27, 28.

Audit Slot	Original Bibliography Reference No.
R1	1
R2	2
R3	3
R4	4
R5	5
R6	6
R7	7
R8	8
R9	9
R10	10

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F - PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: "If I had not verified this citation, how convincing would it look to me?"

Score	Meaning
1	Obviously suspicious
2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	Looks genuine	5	Yes
R2	Looks genuine	5	Yes
R3	Looks genuine	5	Yes
R4	Looks genuine	5	Yes
R5	Looks genuine	5	Yes
R6	Looks genuine	5	Yes
R7	Looks genuine	5	Yes
R8	Looks genuine	5	Yes
R9	Looks genuine	5	Yes
R10	Looks genuine	5	Yes

Example: A citation with realistic authors, a well-known journal, complete volume/pages and a DOI-like identifier may look highly convincing. You might record Plausibility = 5 and Predict Genuine = Yes. You are recording your pre-verification belief, not the truth.

Critical rule: Part F should be completed before Part G. Because this worksheet is being completed after the source-generation step, the entries above are a reconstruction of the initial impression from the unverified bibliography. Replace them with your actual pre-verification answers if you recorded them.

PART G - CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

Persistent identifier: DOI / arXiv / PMID

Publisher or official repository

Google Scholar

Semantic Scholar / OpenAlex

Crossref

PubMed or another discipline-specific database

Exact-title search

Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.
B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves to a different publication.

G4. Examples

A: AI gives a paper and all details match the publisher record.

B: The paper exists, but the AI gives the wrong publication year or journal.

C: The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.

D: You cannot find the paper or a credible scholarly record after systematic searching.

E: The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
R1	Y	Y	Y	Y	Y	Y
R2	Y	Y	Y	Y	Y	Y
R3	Y	Y	Y	Y	Y	Y
R4	Y	Y	Y	Y	Y	Y
R5	Y	Y	Y	Y	Y	Y
R6	Y	Y	Y	Y	Y	Y
R7	Y	Y	Y	Y	Y	Y
R8	Y	Y	Y	Y	Y	Y
R9	Y	Y	Y	Y	Y	Y
R10	Y	Y	Y	Y	Y	Y

Ref.	Verification Source(s)	Identifier Checked	What You Found
R1	arXiv 2210.03629 / DOI 10.48550/arXiv.2210.03629	arXiv 2210.03629	Publication exists; title, authors, year and identifier matched.
R2	NeurIPS proceedings / DOI 10.52202/075280-2997	10.52202/075280-2997	Publication exists; title, authors, venue, year and DOI matched.
R3	Microsoft Research / arXiv 2308.08155	arXiv 2308.08155	Publication exists; title, authors, year and arXiv identifier matched.
R4	arXiv 2303.11366	arXiv 2303.11366	Publication exists; title, authors, year and identifier matched.
R5	arXiv 2112.09332	arXiv 2112.09332	Publication exists; title, authors, year and identifier matched.
R6	arXiv 2308.03688	arXiv 2308.03688	Publication exists; title, authors, year and identifier matched.
R7	arXiv 2303.17651	arXiv 2303.17651	Publication exists; title, authors, year and identifier matched.
R8	arXiv 2307.03172	arXiv 2307.03172	Publication exists; title, authors, year and identifier matched.
R9	arXiv 2405.15793 / Princeton publication record	arXiv 2405.15793	Publication exists; title, authors, year and identifier matched.
R10	arXiv 2408.06292	arXiv 2408.06292	Publication exists; title, authors, year and identifier matched.

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A	Publication exists and the main bibliographic details and identifier match.
R2	A	Publication exists and the main bibliographic details and identifier match.

R3	A	Publication exists and the main bibliographic details and identifier match.
R4	A	Publication exists and the main bibliographic details and identifier match.
R5	A	Publication exists and the main bibliographic details and identifier match.
R6	A	Publication exists and the main bibliographic details and identifier match.
R7	A	Publication exists and the main bibliographic details and identifier match.
R8	A	Publication exists and the main bibliographic details and identifier match.
R9	A	Publication exists and the main bibliographic details and identifier match.
R10	A	Publication exists and the main bibliographic details and identifier match.

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G - SCORING

Count your final classifications: A = 10 B = 0 C = 0 D = 0 E = 0

Total audited: N = A + B + C + D + E = 10

Code	Points
A	4
B	3
C	1
D	0
E	1

Authenticity Points = $4A + 3B + C + E$

Authenticity Score = $[(4A + 3B + C + E) / (4N)] \times 100$

My calculation: $4(10) + 3(0) + 1(0) + 0(0) + 1(0) = 40$ points

Authenticity Score = $40 / (4 \times 10) \times 100 = 100/100$

Worked example: If A=6, B=2, C=1, D=1, E=0, then points = $4(6)+3(2)+1(1)=31$; maximum = 40; Authenticity Score = $31/40 \times 100 = 77.5$.

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59	High citation-authenticity risk
Below 40	Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

Prediction Accuracy = (Correct genuine/fake predictions / Total predictions) x 100

Correct predictions = 10 Total predictions = 10 Prediction Accuracy = 100 %

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	10
References deeply audited	10
A: Verified	10
B: Wrong metadata	0
C: Frankenstein	0
D: Fabricated	0
E: Identifier mismatch	0
Authenticity Score	100 /100
Prediction Accuracy	100 %

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

No. In this audit, all 10 selected references turned out to be genuine, but they looked convincing before verification. The important point is that their appearance was not enough; each one still had to be checked against a scholarly or publisher record.

2. Did all genuine references actually support the claims for which they were cited?

For the main claims checked in the paper, the genuine references supported the claims for which they were cited. No major claim-citation mismatch was found in the audited sample.

3. What was the most serious citation-integrity problem you observed?

No major fabrication, wrong-metadata case, or identifier mismatch was found in the 10 audited references. The main risk was that a convincing-looking citation could easily be trusted without checking it.

4. What is the single most important lesson you learned from this lab?

The biggest lesson is that a citation should not be trusted just because it looks correct. We need to check both whether the source exists and whether it actually supports the claim.

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name

Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name: ...Chintan Maneesh Gangurde.....

Roll Number: ..MHM2026003.....

Signature: ..Chintan Maneesh Gangurde.....

Date: ..21 August 2026.....

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support
Verification -> Quantitative Scores -> Submission